

Engineering an Efficient Probabilistic Exact Model Counter

Mate Soos, Kuldeep S. Meel

University of Toronto, Georgia Tech

24th of July 2025, Zagreb, CAV 2025

Outline

- 1 Propositional Model Counting
- 2 Our Enhancements
- 3 Results

Propositional Model Counting

- Count the number of solutions to a CNF, e.g. $(x_1 \vee x_2) \wedge (\neg x_1 \vee x_3)$
- Notice that finding a single solution is in NP. Counting is #P
- Sometimes, we project the solution space, like so:

Projection set	Other variables


0001	010111
0001	101010
1011	111010

This problem only has 2 solutions over the projection set (but 3 overall)

- Sometimes, literals in the projection set have a “weight”
- Weight is usually a rational number, but it can be any element of a semiring
- Then, each solution has a “weight”: the product of each literal’s weight
- Final solution in this case is the sum of each solution’s weight, i.e. **sum-of-products**

Our Contributions to d-DNNF Model Counting

We count solutions to the CNF by compiling the CNF into a **d-DNNF** [Darwiche&Marquis, 2002]: a deterministic decomposable negation normal form circuit.



- ① **Enhanced Residual Formula Processing:** optimized SAT solver architecture for residual formula processing that incorporates VSIDS scoring, restarts, and polarity caching.
- ② **Dual Independent Set Framework:** maintains distinct SAT-eligibility (S) and decision (D) sets, where the S-set determines SAT solver transitions while the D-set guides branching decisions.
- ③ **Chronological Backtracking:** Adaptation of chronological backtracking to model counting


Calling a SAT solver

Once all variables in P have been decided there is only one solution: we can run a SAT solver!

$$V = \{x_1, x_2, x_3, x_4, \dots\}$$

$$P = \{x_1, x_2\}$$

Once $\{x_1, x_2\}$ have all been branched on, we can call a SAT solver on the residual formula!



This has been exploited before by gpmc

Enhanced Residual Formula Processing

We added:

- **VSIDS Scoring**: Model counters use VSADS¹ scoring, but for SAT, VSIDS is faster to compute, and more efficient.
- **Restarts**: Restart the SAT solver regularly (Luby heuristic²), to explore different parts of the search space. Note we only need one solution.
- **Polarity Caching**³: We cache the polarity of variables to avoid unnecessary recomputation, especially due to restarts

Basically: let's lift all the SAT techniques to model counting's residual formula processing.

¹Variable State Aware Decaying Sum, from Sang et al.'05

²Luby et al, '93

³Pipatsrisawat et al, '07


Dual Independent Set

Definition (Dual Independent Set)

Given a Boolean formula F defined over variables V and a projection set \mathcal{P} , a dual independent set consists of:

- ① **S-set** A SAT-eligibility set $S \subseteq \mathcal{P}$ that determines when SAT solver mode transition is permissible.
- ② **D-set** A decision set $D \subseteq V$ that guides branching variable selection, where $D \supseteq \mathcal{P}$.

This formulation generalizes traditional approaches where $D = S \subseteq P$, instead aiming for $S \subseteq P \subseteq D$.



Computing Maximal Decision Set

```

1:  $D \leftarrow \mathcal{P}$ ,  $G \leftarrow \text{EXTRACTGATES}(F)$ 
2: procedure SYNTACTICEXPANSION( $V$ )
3:   changed  $\leftarrow V$ 
4:   while changed is not empty do
5:      $v \leftarrow \text{changed.pop}()$ 
6:     for gate  $g \in G$  where  $v$  is input do If possible, extend  $D$  with output of the gate
7:   return  $D$ 
8: end procedure
9: procedure SEMANTICEXPANSION
10:  for  $v \notin D$  do
11:    if VALIDATEDECISIONVAR( $D \cup \{v\}$ ,  $F$ ) then
12:       $D \leftarrow D \cup \{v\}$ 
13:       $D \leftarrow \text{SYNTACTICEXPANSION}(\{v\})$ 
14:  return  $D$ 
15: end procedure

```

▷ Initialize D with \mathcal{P} , Extract gates
 ▷ Quick check with syntactic analysis

We run SYNTACTICEXPANSION(P) and then SEMANTICEXPANSION

Chronological Backtracking

Notice that we should NOT backjump or we lose already counted solutions.

- Model counters go back to the deepest level possible, flip the decision, and throw away the learnt clause in case it would violate propagation invariants
- But we can relax some of the invariants: Chronological Backtracking (ChronoBT)
- ChronoBT allows the have out-of-order implication levels in the trail
- ChronoBT allows us to keep the learnt clause, and force the literal that the learnt clause entails.
- We used the fuzzer SharpVelvet by Latour et al. to find bugs in our

$$c_1 = (\neg V_1 \vee V_2)$$

$$c_2 = (\neg V_4 \vee V_5)$$

$$c_3 = (\neg V_5 \vee V_6)$$

$$c_4 = (\neg V_2 \vee \neg V_4 \vee \neg V_6)$$


$$c_5 = (\neg V_3 \vee V_7)$$

$$c_6 = (\neg V_7 \vee \neg V_4 \vee V_8)$$

$$c_7 = (V_4 \vee V_9)$$

$$c_8 = (\neg V_1 \vee V_4 \vee \neg V_9)$$

Input Formula



Conflict 1 : Implication Graph

Conflicting clause: c_4
1UIP conflict clause:

$$c_9 = (\neg V_2 \vee \neg V_4)$$

$$d_3 \neg V_6 \quad \text{X}$$

$$d_3 V_6$$

$$d_3 V_5$$

$$d_3 V_4$$

$$d_2 V_7$$

$$d_2 V_3$$

$$d_1 V_2$$


$$d_1 V_1$$

Conflict 1 : Trail

Figure from [Chronological Backtracking by Nadel&Ryvchin, 2018](#)


Compensating Weights

- ① We explore the left side of a graph, nodes 1 and 2
- ② We backtrack to level 0 and explore the right side, nodes 3 and 4.
- ③ At node 4, we find the unit clause x_3 . This unit clause's level is 0, but due to ChronoBT, we only backtrack to level 1. However, the system **has already multiplied in the weight of x_3** into nodes 1, 2, and 3.
- ④ These nodes' weights, which on the left side of an already explored branch need to be compensated



Compensating Weights

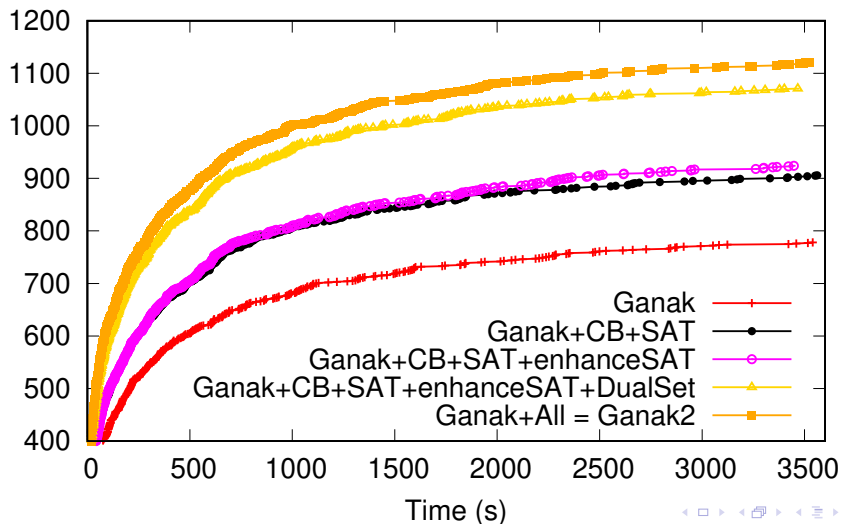
- ① x_4 is part of components $\{1,2,3,4\}$ but not 5 is learned to be false at level 0.
- ② Sounds impossible. x_4 is clearly not part of component 5 (since it is part of 3 and 4), so it should never be part of a learned clause while examining component 5
- ③ Components are decided *purely* based on irredundant clauses. It is possible that learned clauses connect components!
- ④ Can lead to contradictions over variables *not* part of the component examined!
- ⑤ Compensating for these is non-trivial: *sibling* components need to be examined and compensated for




Compensating Weights

- Notice that to compensate we need to **divide**
- But weighted counting can be done with weights over any **semiring**
- But Ganak2 needs a **field**
- In some sense, we reorder computations, exploiting flexibility in building the d-DNNF
- Side-note: Ganak2 supports **any field**. Implementing a new field should take ≈ 10 minutes. Currently supports:
 - Integers
 - Counting modulo prime
 - Rational numbers
 - Floating point – not a field actually
 - Complex rational numbers
 - Multivariate polynomials over the rationals

Ganak Versions over all Instances



All Instances, All Solvers



Conclusions

We created a new probabilistically exact model counter, Ganak2, that incorporates:

- **Enhanced Residual Formula Processing**
- **Dual Independent Set Framework**
- **Chronological Backtracking**

This allowed us to achieve state-of-the-art performance. Ganak2 won all tracks of the Model Counting Competition 2024.

Code:



Paper:



Also, try it out online: <https://www.msoos.org/ganak/>